

Exercise 1

Bowen Hua

September 16, 2017

1 Linear Regression

1.1 (A)

Matrix form:

$$\hat{\beta} = \arg \min_{\beta \in \mathcal{R}^P} \frac{1}{2} (y - X\beta)^T W (y - X\beta).$$

Since W is symmetric and positive semidefinite, this is a convex optimization problem. Its first-order optimality condition is necessary and sufficient:

$$X^T W (y - X\beta) = \mathbf{0}.$$

1.2 (B)

1.2.1 Method 1: direct inversion

Described in the problem statement:

$$\beta = (X^T W X)^{-1} X^T W y.$$

To exploit sparsity of W , we use broadcasting in Python instead of matrix multiplication to compute operations w.r.t. W matrix. This method has a complexity of $O(np^2)$.

1.2.2 Method 2: pseudoinverse

Since the weight matrix W is diagonal, we can define $W^{\frac{1}{2}}$ to be a diagonal matrix where the i -th diagonal element equals $\sqrt{w_i}$. Now we can re-write the optimality condition:

$$\beta = [(W^{\frac{1}{2}} X)^T (W^{\frac{1}{2}} X)]^{-1} (W^{\frac{1}{2}} X)^T y,$$

which we re-write as

$$\beta = (W^{\frac{1}{2}} X)^{\dagger} W^{\frac{1}{2}} y,$$

where $(W^{\frac{1}{2}} X)^{\dagger}$ is the pseudoinverse of $W^{\frac{1}{2}} X$.

We first “preprocess” the feature matrix X and y by multiplying $W^{\frac{1}{2}}$ on the left.

We then compute the pseudoinverse through computing SVD of $W^{\frac{1}{2}} X$. This method is numerically more stable than the direct inverse method. (There could be correlation between our observations X . Therefore we care about numerical stability.)

Pseudocode for pseudoinverse of matrix A :

```
(U, Σ, V) = svd(A)
for Σii in Σ: #traverse through the diagonal elements
    if Σii ≠ 0:
        Σii = 1/Σii
return VTΣUT
```

In the Python code, we call `numpy.linalg.pinv` to perform the pseudoinverse. This method also has a complexity of $O(np^2)$.

In addition, this is the implementation of `scikit-learn`.

1.2.3 Method 3: Cholesky decomposition

We can use Cholesky decomposition on the matrix $X^T W X$. Now we have $LL^T \beta = D$. Then we can obtain β by solving two linear systems.

This method also has a complexity of $O(np^2)$.

Pseudocode for Cholesky-decomposition-based method:

```
Let  $C = X^T W X$ ,  $d = X^T W y$ .
Compute Cholesky decomposition  $C = LL^T$ .
Solve for  $\alpha$  in  $L\alpha = d$ .
Solve for  $\beta$  in  $L^T \beta = \alpha$ .
```

1.3 (C)

I coded the three methods in Python, using the `numpy` package. The codes can be found in my GitHub.

The results are summarized here:

Table 1: CPU Times (s) for Three Methods of Weighted Least Squares

(n, p)	Method 1	Method 2	Method 3
(2000, 50)	0.001	0.007	0.001
(1000, 1000)	0.122	0.346	0.064
(20000, 50)	0.012	0.034	0.005
(50000, 50)	0.028	0.118	0.013
(5000, 5000)	6.402	37.00	4.462

We can see that:

- Method 3 (Cholesky) consistently performs better than Method 1 (Direct Inverse), which is faster than Method 2 (pseudoinverse).
- When X is close to a square matrix, the performance of Method 3 is way worse than the other two methods.


1.4 (D)

We use `scipy.sparse` in Python as our tool for sparse matrix operations. In particular, we use `scipy.sparse.linalg.lsqr` to solve the sparse least square problem:

$$\beta = (W^{\frac{1}{2}}X)^{\dagger}W^{\frac{1}{2}}y.$$

We generate random X matrices of size $(200000, 50)$, with different density. We name the sparse method as Method 4. The results are summarized as follows:

Figure 1: CPU Times (s) of Four Methods for the Sparse Matrix



We can see that:

- Method 1–3 do not exploit sparsity. Their CPU times do not change with density.
- When density is relatively small, the sparse method has an advantage over all the other three method. When density is relatively large, the sparse method is slower than Method 1 and 3, possibly due to overhead of the sparse data structure.

2 Generalized Linear Models

2.1 (A)

We have $y_i \sim \text{Binomial}(m_i, w_i)$ ¹, where

$$w_i = \frac{1}{1 + \exp(-x_i^T \beta)},$$

$$1 - w_i = \frac{\exp(-x_i^T \beta)}{1 + \exp(-x_i^T \beta)}.$$

The negative log likelihood function is:

$$\ell(\beta) = -\log \left\{ \prod_{i=1}^N p(y_i | \beta) \right\} \quad (1)$$

$$= -\log \left\{ \prod_{i=1}^N \binom{m_i}{y_i} (w_i)^{y_i} (1 - w_i)^{m_i - y_i} \right\} \quad (2)$$

$$= - \left\{ \sum_{i=1}^N \left(\log \binom{m_i}{y_i} + y_i \log(w_i) + (m_i - y_i) \log(1 - w_i) \right) \right\} \quad (3)$$

We have:

$$\nabla \log w_i = -\nabla \log(1 + \exp(-x_i^T \beta)) \quad (4)$$

$$= \frac{x_i \exp(-x_i^T \beta)}{1 + \exp(-x_i^T \beta)} \quad (5)$$

$$= x_i(1 - w_i) \quad (6)$$

$$\nabla \log(1 - w_i) = \nabla(-x_i^T \beta) - \nabla \log(1 + \exp(-x_i^T \beta)) \quad (7)$$

$$= -x_i + x_i(1 - w_i) \quad (8)$$

$$= -x_i w_i \quad (9)$$

Therefore, the gradient of the loss function is:

$$\nabla \ell(\beta) = \sum_{i=1}^N (y_i x_i (1 - w_i) - (m_i - y_i) x_i w_i) \quad (10)$$

$$= \sum_{i=1}^N (m_i w_i - y_i) x_i \quad (11)$$

$$= X^T (m \circ w - y), \quad (12)$$

where operator \circ denotes element-wise product.

¹For the simpler case of binary logistic regression where we have $y_i \sim \text{Bernoulli}(w_i)$, just apply $m_i = 1$ in the following results.

2.2 (B)

We use the data `wdbc.csv` from course website. We use the first 10 features in our model. We add a column of ones into X matrix to represent the intercept term.


To alleviate the numerical issue brought by having a very large w_i value, we first scale the X data using `scikit-learn` `preprocessor`.

We perform a naive gradient descent with different step lengths:

- fixed step size of 0.001,
- fixed step size of 0.005,
- fixed step size of 0.01, and
- variable step size of $1/k$ where k is the iteration count.

To clearly see the difference between different step sizes, we compute 50 iterations. The codes are shown in `logit.py` and the loss function value as a function of iteration count is shown below:

Figure 2: Loss function value as a function of iteration count



The step size 0.01 performs well for this dataset. For the first few iterations, the step sizes for the $1/k$ rule are too large, so that the loss function does not decrease as much.

I also implemented a `predict()` function that evaluates the training error of our logistic model. After only 50 iterations, the training accuracy is 94.02%.

2.3 (C)

Our objective is to compute the Taylor approximation:

$$\hat{\ell}(\beta) = \ell(\beta_0) + (\nabla \ell(\beta))^T (\beta - \beta_0) + \frac{1}{2} (\beta - \beta_0)^T \nabla^2 \ell(\beta) (\beta - \beta_0).$$

We have computed $\nabla \ell(\beta)$ in section (A). Now we compute the (i, j) element of the Hessian $\nabla^2 \ell(\beta)$:

$$\frac{\partial^2}{\partial \beta_i \partial \beta_j} \ell(\beta) = \frac{\partial}{\partial \beta_i} (\nabla_j \ell(\beta)) \quad (13)$$

$$= \frac{\partial}{\partial \beta_i} \left(\sum_{k=1}^N (m_k w_k - y_k) x_{kj} \right) \quad (14)$$

$$= \sum_{k=1}^N x_{ki} x_{kj} m_k w_k (1 - w_k), \quad (15)$$

where we use the fact that

$$\frac{\partial}{\partial \beta_i} w_k = x_{ki} w_k (1 - w_k).$$

We can define a new diagonal matrix $W = \text{diag}(m_1 w_1 (1 - w_1), \dots, m_N w_N (1 - w_N))$, then we can write the Hessian in matrix form:

$$\nabla^2 \ell(\beta) = X^T W X.$$

Plugging in the expressions of the gradient and Hessian into the Taylor approximation yields:

$$\hat{\ell}(\beta) = \ell(\beta_0) + (X^T (m \circ w - y))^T (\beta - \beta_0) + \frac{1}{2} (\beta - \beta_0)^T X^T W X (\beta - \beta_0).$$

Since we do not care about the constant term, we do not keep track of it in the following derivation. Instead, we use c to represent the constant term, without specifying what c is. Using the technique from: <https://justindomke.wordpress.com/completing-the-square-in-n-dimensions/>, we can continue our reformation:

$$\begin{aligned} \hat{\ell}(\beta) &= \frac{1}{2} ([\beta - \beta_0] + (X^T W X)^{-1} X^T (m \circ w - y))^T X^T W X ([\beta - \beta_0] + \\ &\quad (X^T W X)^{-1} X^T (m \circ w - y)) + c' \\ &= \frac{1}{2} (\beta - \beta_0 - X^{-1} W^{-1} (m \circ w - y))^T X^T W X (\beta - \beta_0 - X^{-1} W^{-1} (m \circ w - y)) + c' \\ &= \frac{1}{2} (X\beta - X\beta_0 + W^{-1} (m \circ w - y))^T W (X\beta - X\beta_0 + W^{-1} (m \circ w - y)) + c'' \\ &= \frac{1}{2} (z - X\beta)^T W (z - X\beta) + c, \end{aligned}$$

where $z = X\beta_0 + W^{-1} (y - m \circ w)$.

2.4 (D)

Now we use Newton's method to perform our loss minimization. The update rule is as follows:

$$\hat{\beta}_{t+1} = \hat{\beta}_t - (\nabla^2 \ell(\hat{\beta}_t))^{-1} \nabla \ell(\hat{\beta}_t),$$

where we have computed the Hessian in the previous subsection.